


Red-team reconciliation — the five attack paths and where each is stopped

Semantic pattern: **adversarial policy evaluation** → visual type: **flowchart (attack lanes)** · external report RT-01–RT-05 reconciled at HEAD 203e2c2 · 1325 tests



Provenance manifest

diagram source	commit 203e2c2; reconciliation record hermes_redteam_reconciliation.md; external report RT-01–RT-05
files inspected	src/hermes/research/controller.py (record_source_retraction, retracted_source_review_candidates), src/hermes/persistence/source_outcomes.py (PS-08 / HZ-02), docs/idr/IDR-040.md (authority boundary), hermes_current_epistemic_architecture_review.md (§19, §20)

tests pinning each lane	RT-01: test_two_observers_same_retraction_two_audit_events_one_seed, test_forged_retraction_events_never_touch_ladder_state, test_identical_retry_does_not_duplicate_the_event, test_fetch_without_retraction_hazard_emits_no_event · RT-02: test_scope_routes_are_pending, test_non_authority_actions_stay_effective · RT-03: test_goodhart_padding_satisfied_obligations_never_wins, test_goodhart_padding_outstanding_caps_at_high · RT-04: test_succeeded_intermediate_breaks_the_blocking_chain · RT-05: tests/test_provider_hazards.py
probe evidence	two cross-task observers → 2 audit events (distinct observed_by_task), advisory deduped to one seed; forged + corrupt + foreign events → ladder identical, APPLY wrote nothing, 2 fail-closed notes
projection rule	diagrams are never authority; executable code and ratified IDRs win on any conflict (hermes_diagram_design_skill.md)